

# WebAssembly

---

By Jonathan Easterman and Vivek Pradhan

# The Problem

- The browser is the largest platform for applications
  - Statistic or graphic here
  - Electron
  - Google Drive
  - Computer Vision and Machine Learning (Tensorflow.js)
- Javascript is limited
  - Dynamically Typed
  - Automatic Memory Management
  - Single threaded

# Solution - Asm.js

- Precursor to WebAssembly
- Transpile C/C++ to optimized dialect of Javascript
  - Type coercion
  - Interact with memory via typed array
  - enables ahead-of-time compilation
- Results in large .js files
  - Need to include entire JS ports of C libraries
- Limited to expressiveness of Javascript
  - E.g. type expressiveness
- Not standardized between browser vendors

# Asm.js Example

```
size_t strlen(char *ptr) {  
    char *curr = ptr;  
    while (*curr != 0) {  
        curr++;  
    }  
    return (curr - ptr);  
}
```

```
function strlen(ptr) {  
    ptr = ptr|0;  
    var curr = 0;  
    curr = ptr;  
    while (MEM8[curr]|0 != 0) {  
        curr = (curr + 1)|0;  
    }  
    return (curr - ptr)|0;  
}
```

# Enter WebAssembly

- Similar to JVM bytecode
  - Typed → Faster to execute
  - Stack-machine bytecode → Compact → Faster to parse + transmit
- Not a replacement for Javascript
- Currently intended for manually managed languages
  - C/C++ - EmScripten + LLVM
  - Rust - rustc + LLVM

### C input source

```
int factorial(int n) {  
    if (n == 0)  
        return 1;  
    else  
        return n * factorial(n-1);  
}
```

### Linear assembly bytecode (intermediate representation)

```
get_local 0  
i64.eqz  
if (result i64)  
    i64.const 1  
else  
    get_local 0  
    get_local 0  
    i64.const 1  
    i64.sub  
    call 0  
    i64.mul  
end
```

### Wasm binary encoding (hexadecimal bytes)

```
20 00  
50  
04 7E  
42 01  
05  
20 00  
20 00  
42 01  
7D  
10 00  
7E  
0B
```

# Details

- Four types
  - 32 + 64-bit variants of int + float
- `WebAssembly.Module()`
  - Stateless Native code
- `WebAssembly.Instance()`
  - Stateful Instantiated module
- `WebAssembly.Memory()`
  - Resizable `ArrayBuffer` of raw bytes
  - Like a Heap for an Instance
- `WebAssembly.Table()`
  - Resizable typed array of opaque values
  - Currently used for function pointers

# Hypotheses

Goal: Compare JS, Asm.js, WebAssembly



- WebAssembly will perform better for computationally intensive tasks
- WebAssembly will suffer for tasks that involve moving/copying from Linear Memory to JS memory



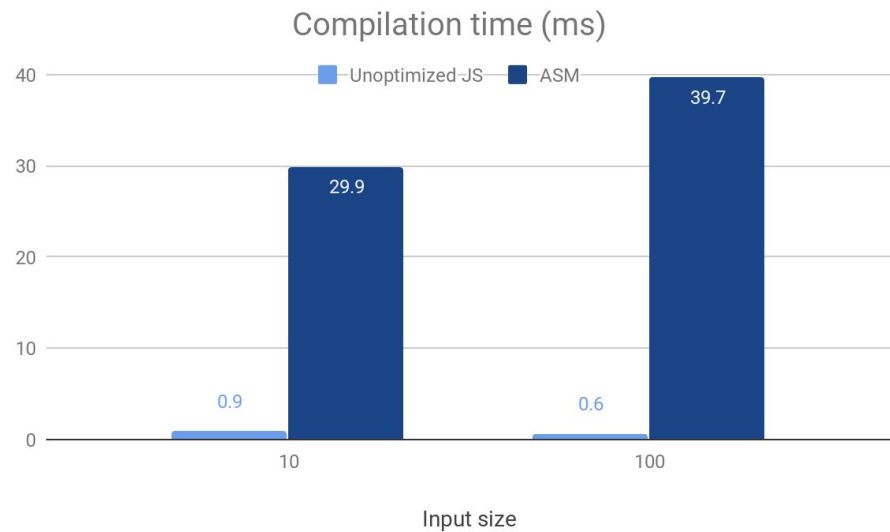
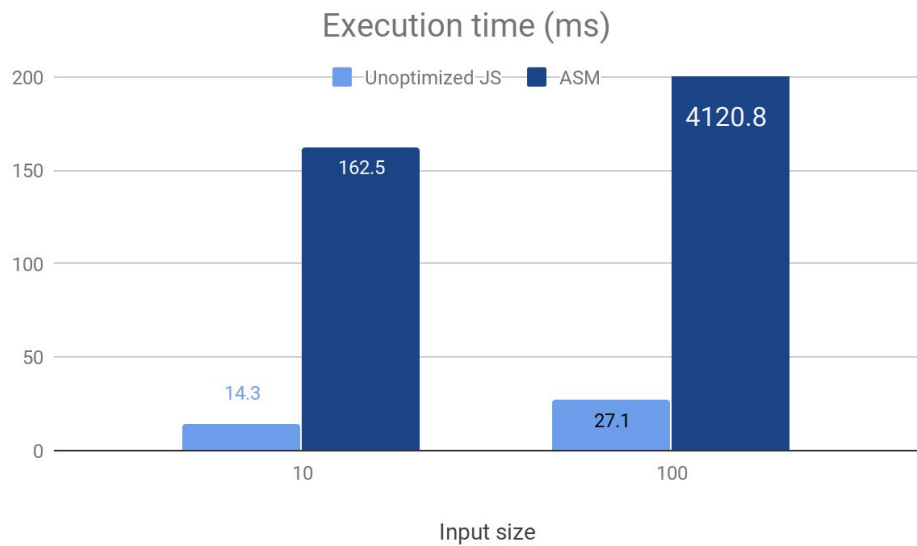
# Experimental Setup

- 3 programs written in JS and C
  - monte-carlo
  - spectral-norm
  - deepcopy
- EmScripten compiler
  - C → LLVM → WebAssembly
  - C → LLVM → Asm.js
  - Includes implementation of C stdlib
- Collect statistics via V8 profiler and Chrome Dev tools
  - Sample size:  $n = 10$

# Monte Carlo

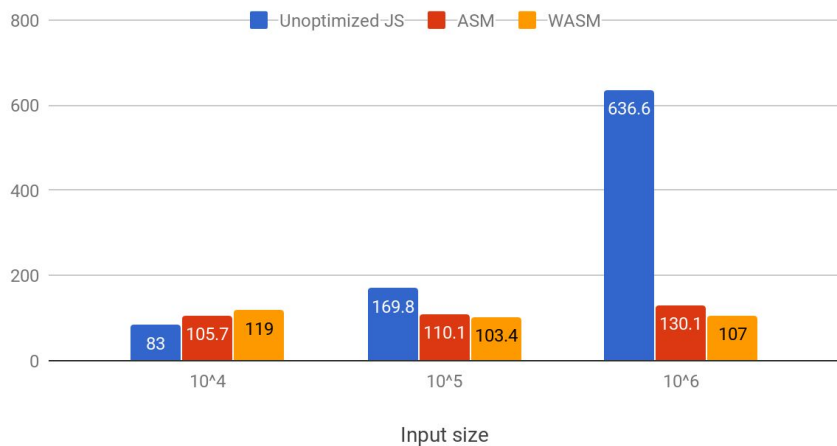


# Spectral Norm




# Deep copy

Execution time (ms)



Compilation time (ms)



# WebAssembly Limitations

- Does not support SIMD
- No GC
- Relies on JS APIs for many things
  - E.g. DOM manipulation
- Single threaded
- Early support in browsers
- Tooling is not the most accessible

# Conclusions

- WebAssembly is a promising tool for stealing performance back
- Opens web development to performance-oriented systems engineers
- In practice, hypothesize and test!!